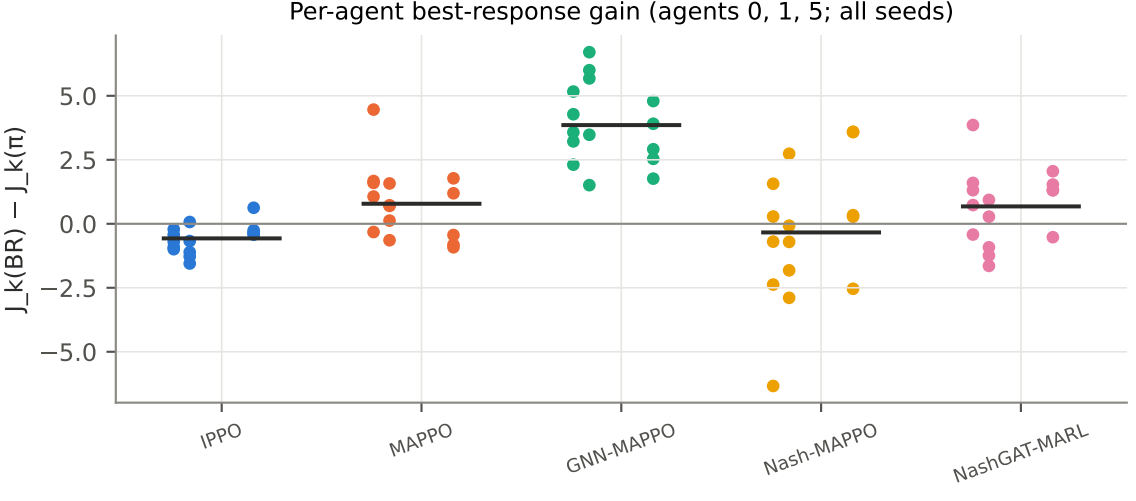
Per-agent best-response gain (agents 0, 1, 5; all seeds)
J_k(BR) − J_k(π)
5.0
2.5
0.0
−2.5
−5.0
IPPO
MAPPO
GNN-MAPPO
Nash-MAPPO
NashGAT-MARL